Dataset
- Around 200k ratings and 5k books
- Ratings are saved when they are made
- All in CSV files

User profiling
- Stores users forename, surname and username
- Also, store the localisation that the user selected
- Allows the user to store many ratings of books

Recommender
- SVD calculated each time a new user is added and otherwise would be periodic, this would allow for the use of a much bigger dataset
- Use matrices U and V from SVD to calculate each user's latent factors and compare them to books in manhattan distance
- This means that we can rank the books in how close they are to the user without doing SVD each time the user makes a new rating, giving "instant" responses no matter the size of the dataset
- This gives quite interesting recommendations, although it seems to do better on ones that have more ratings in my dataset, and some with few ratings get shown too frequently

Interface
- Uses bootstrap for responsive HTML
- Jquery to manage ajax and data handling on dynamic pages
- Allows the user to change the language persistently in settings (only while logged in)